

# Reproducible Research: Peer Assessment 1

*William L. Kennedy*

## Loading and preprocessing the data

```
# Package dependencies (uncomment and install if necessary)

# install.packages("data.table")
# install.packages("lubridate")
library(data.table)
library(lubridate)

##
## Attaching package: 'lubridate'
##
## The following objects are masked from 'package:data.table':
##
##     hour, mday, month, quarter, wday, week, yday, year

# set the working directory (replace path accordingly)
setwd("/Users/adakemia/Documents/Academic/Coursera/DataScienceSpecialization/05ReproducibleResearch/Pro

# List and check for necessary files
url <- "https://d396qusza40orc.cloudfront.net/repdata%2Fdata%2Factivity.zip"
zfile <- "activity.zip"
file <- "activity.csv"

if (!file.exists(file)) {
  if (!file.exists(zfile)) {
    method <- switch(Sys.info()[['sysname']],
                     "Windows" = "internal",
                     "Darwin" = "curl",
                     "Linux" = "wget",
                     "auto")
    download.file(url, destfile=zfile, method)
  }
  unzip(zfile, exdir=".")
}

# Read files into data.table object
data <- fread(file, sep=";", stringsAsFactors=F, header=T,
              na.strings="NA", colClasses=c("numeric","character","character"))
# Check structure of data
str(data)

## Classes 'data.table' and 'data.frame':  17568 obs. of  3 variables:
## $ steps   : num  NA NA NA NA NA NA NA NA NA NA NA ...
## $ date    : chr  "2012-10-01" "2012-10-01" "2012-10-01" "2012-10-01" ...
## $ interval: chr  "0" "5" "10" "15" ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
# Process data
```

```
# Set date formatting (lubridate package)  
data[,date := ymd(date)]
```

```
##      steps      date interval  
##    1:    NA 2012-10-01         0  
##    2:    NA 2012-10-01         5  
##    3:    NA 2012-10-01        10  
##    4:    NA 2012-10-01        15  
##    5:    NA 2012-10-01        20  
##    ---  
## 17564:    NA 2012-11-30       2335  
## 17565:    NA 2012-11-30       2340  
## 17566:    NA 2012-11-30       2345  
## 17567:    NA 2012-11-30       2350  
## 17568:    NA 2012-11-30       2355
```

```
# Data quality checks
```

```
# Check for missingness  
colSums(is.na(data))
```

```
##      steps      date interval  
##    2304         0           0
```

```
# Percent missing
```

```
sum(is.na(data$steps)) / nrow(data) * 100
```


```
## [1] 13.11475
```

```
# Check timeseries
```

```
data[,unique(date)]
```

```
## [1] "2012-10-01 UTC" "2012-10-02 UTC" "2012-10-03 UTC" "2012-10-04 UTC"  
## [5] "2012-10-05 UTC" "2012-10-06 UTC" "2012-10-07 UTC" "2012-10-08 UTC"  
## [9] "2012-10-09 UTC" "2012-10-10 UTC" "2012-10-11 UTC" "2012-10-12 UTC"  
## [13] "2012-10-13 UTC" "2012-10-14 UTC" "2012-10-15 UTC" "2012-10-16 UTC"  
## [17] "2012-10-17 UTC" "2012-10-18 UTC" "2012-10-19 UTC" "2012-10-20 UTC"  
## [21] "2012-10-21 UTC" "2012-10-22 UTC" "2012-10-23 UTC" "2012-10-24 UTC"  
## [25] "2012-10-25 UTC" "2012-10-26 UTC" "2012-10-27 UTC" "2012-10-28 UTC"  
## [29] "2012-10-29 UTC" "2012-10-30 UTC" "2012-10-31 UTC" "2012-11-01 UTC"  
## [33] "2012-11-02 UTC" "2012-11-03 UTC" "2012-11-04 UTC" "2012-11-05 UTC"  
## [37] "2012-11-06 UTC" "2012-11-07 UTC" "2012-11-08 UTC" "2012-11-09 UTC"  
## [41] "2012-11-10 UTC" "2012-11-11 UTC" "2012-11-12 UTC" "2012-11-13 UTC"  
## [45] "2012-11-14 UTC" "2012-11-15 UTC" "2012-11-16 UTC" "2012-11-17 UTC"  
## [49] "2012-11-18 UTC" "2012-11-19 UTC" "2012-11-20 UTC" "2012-11-21 UTC"  
## [53] "2012-11-22 UTC" "2012-11-23 UTC" "2012-11-24 UTC" "2012-11-25 UTC"  
## [57] "2012-11-26 UTC" "2012-11-27 UTC" "2012-11-28 UTC" "2012-11-29 UTC"  
## [61] "2012-11-30 UTC"
```

```
plot(data$interval[11:14], type="l")
```



```
data[11:14]
```

```
##      steps      date interval
## 1:    NA 2012-10-01      50
## 2:    NA 2012-10-01      55
## 3:    NA 2012-10-01     100
## 4:    NA 2012-10-01     105
```

```
# Fix hour/min anomaly
```

## What is mean total number of steps taken per day?

First, we can take a wide view via a histogram. We can see several things from the histogram:


The shape is relatively normal

Due to normality, we would expect mean and median to be fairly close

We would expect the mean and median to fall roughly just over 10,000 steps

```
hist(data[,sum(steps), by=date]$V1, breaks=8,
      main="Histogram of Mean Steps Per Day",
      xlab="mean steps per day")
```

## Histogram of Mean Steps Per Day



Next, we can compare to the actual mean and median for the entire period. As expected, the mean and median are very close and just over 10,000 steps.

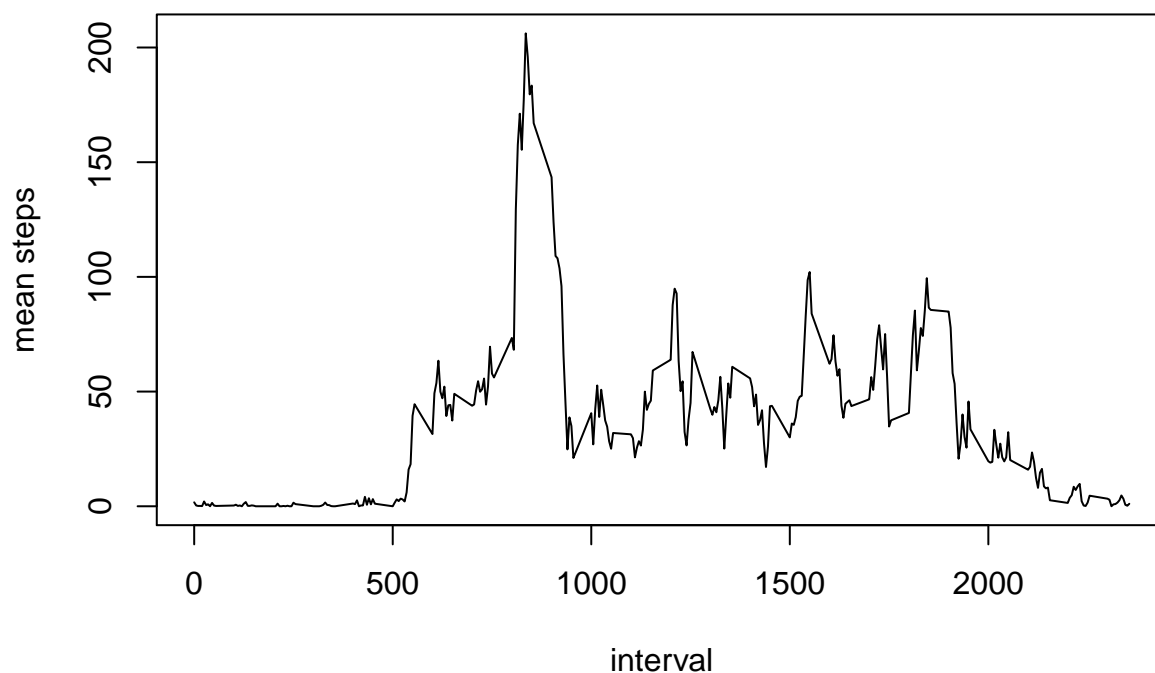
```
data[, .(sum = sum(steps)), by=date][, .(median = median(sum, na.rm=T),  
                                         mean = mean(sum, na.rm=T))]
```

```
##      median      mean  
## 1:  10765 10766.19
```

What is the average daily activity pattern?

```
ts <- data[, .(mean = round( mean(steps, na.rm=T), 2)), by=interval]  
  
plot(ts$mean ~ ts$interval, type="l", main="Mean Steps Per Daily Interval",  
     xlab="interval", ylab="mean steps")
```

## Mean Steps Per Daily Interval



```
ts[which.max(mean)]
```


```
##      interval  mean
## 1:         835 206.17
```

## Imputing missing values

```
# create a second imputed data set
data_imp <- data

ts2 <- data_imp[, .(median = median(steps, na.rm=T),
                    mean = mean(steps, na.rm=T)),
                 by=interval]
# Compare mean and median by interval
plot(ts2$mean ~ ts2$interval, type="l", main="Mean Steps Per Daily Interval",
     xlab="interval", ylab="mean steps")
lines(ts2$median ~ ts2$interval, col="red")
```

## Mean Steps Per Daily Interval



```
data[interval == 1000]
```

##	steps	date	interval
## 1:	NA	2012-10-01	1000
## 2:	0	2012-10-02	1000
## 3:	0	2012-10-03	1000
## 4:	0	2012-10-04	1000
## 5:	0	2012-10-05	1000
## 6:	16	2012-10-06	1000
## 7:	281	2012-10-07	1000
## 8:	NA	2012-10-08	1000
## 9:	23	2012-10-09	1000
## 10:	400	2012-10-10	1000
## 11:	0	2012-10-11	1000
## 12:	0	2012-10-12	1000
## 13:	0	2012-10-13	1000
## 14:	392	2012-10-14	1000
## 15:	0	2012-10-15	1000
## 16:	0	2012-10-16	1000
## 17:	92	2012-10-17	1000
## 18:	0	2012-10-18	1000
## 19:	0	2012-10-19	1000
## 20:	0	2012-10-20	1000
## 21:	0	2012-10-21	1000
## 22:	0	2012-10-22	1000
## 23:	0	2012-10-23	1000
## 24:	0	2012-10-24	1000
## 25:	0	2012-10-25	1000
## 26:	0	2012-10-26	1000

```
## 27:      0 2012-10-27      1000
## 28:      0 2012-10-28      1000
## 29:    104 2012-10-29      1000
## 30:      0 2012-10-30      1000
## 31:    122 2012-10-31      1000
## 32:     NA 2012-11-01      1000
## 33:    487 2012-11-02      1000
## 34:      0 2012-11-03      1000
## 35:     NA 2012-11-04      1000
## 36:      0 2012-11-05      1000
## 37:      0 2012-11-06      1000
## 38:      8 2012-11-07      1000
## 39:      0 2012-11-08      1000
## 40:     NA 2012-11-09      1000
## 41:     NA 2012-11-10      1000
## 42:      0 2012-11-11      1000
## 43:      0 2012-11-12      1000
## 44:      0 2012-11-13      1000
## 45:     NA 2012-11-14      1000
## 46:      0 2012-11-15      1000
## 47:     18 2012-11-16      1000
## 48:      0 2012-11-17      1000
## 49:      0 2012-11-18      1000
## 50:      0 2012-11-19      1000
## 51:      0 2012-11-20      1000
## 52:      0 2012-11-21      1000
## 53:      0 2012-11-22      1000
## 54:      0 2012-11-23      1000
## 55:      0 2012-11-24      1000
## 56:      0 2012-11-25      1000
## 57:      0 2012-11-26      1000
## 58:    207 2012-11-27      1000
## 59:      0 2012-11-28      1000
## 60:      0 2012-11-29      1000
## 61:     NA 2012-11-30      1000
##      steps      date interval
```

```
data_imp[, .(median = median(steps, na.rm=T),
              mean = mean(steps, na.rm=T)),
           by=list(interval, wday(date))][interval==900]
```

```
##      interval wday median      mean
## 1:         900    2      71 249.85714
## 2:         900    3      36 134.00000
## 3:         900    4      14 137.50000
## 4:         900    5       0 137.75000
## 5:         900    6      16 218.14286
## 6:         900    7      20 112.14286
## 7:         900    1      15  19.14286
```

```
ts2[interval == 900]
```

```
##      interval median      mean
## 1:         900      20 143.4528
```

```
data[interval == 900]
```

##	steps	date	interval
## 1:	NA	2012-10-01	900
## 2:	0	2012-10-02	900
## 3:	0	2012-10-03	900
## 4:	0	2012-10-04	900
## 5:	530	2012-10-05	900
## 6:	30	2012-10-06	900
## 7:	23	2012-10-07	900
## 8:	NA	2012-10-08	900
## 9:	134	2012-10-09	900
## 10:	135	2012-10-10	900
## 11:	548	2012-10-11	900
## 12:	802	2012-10-12	900
## 13:	18	2012-10-13	900
## 14:	15	2012-10-14	900
## 15:	732	2012-10-15	900
## 16:	36	2012-10-16	900
## 17:	668	2012-10-17	900
## 18:	55	2012-10-18	900
## 19:	15	2012-10-19	900
## 20:	0	2012-10-20	900
## 21:	0	2012-10-21	900
## 22:	71	2012-10-22	900
## 23:	400	2012-10-23	900
## 24:	0	2012-10-24	900
## 25:	0	2012-10-25	900
## 26:	164	2012-10-26	900
## 27:	519	2012-10-27	900
## 28:	23	2012-10-28	900
## 29:	4	2012-10-29	900
## 30:	0	2012-10-30	900
## 31:	0	2012-10-31	900
## 32:	NA	2012-11-01	900
## 33:	16	2012-11-02	900
## 34:	198	2012-11-03	900
## 35:	NA	2012-11-04	900
## 36:	363	2012-11-05	900
## 37:	433	2012-11-06	900
## 38:	28	2012-11-07	900
## 39:	0	2012-11-08	900
## 40:	NA	2012-11-09	900
## 41:	NA	2012-11-10	900
## 42:	0	2012-11-11	900
## 43:	539	2012-11-12	900
## 44:	0	2012-11-13	900
## 45:	NA	2012-11-14	900
## 46:	0	2012-11-15	900
## 47:	0	2012-11-16	900
## 48:	20	2012-11-17	900
## 49:	11	2012-11-18	900
## 50:	0	2012-11-19	900



```
## 51: 203 2012-11-20 900
## 52: 269 2012-11-21 900
## 53: 499 2012-11-22 900
## 54: 0 2012-11-23 900
## 55: 0 2012-11-24 900
## 56: 62 2012-11-25 900
## 57: 40 2012-11-26 900
## 58: 0 2012-11-27 900
## 59: 0 2012-11-28 900
## 60: 0 2012-11-29 900
## 61: NA 2012-11-30 900
##      steps      date interval
```

**Are there differences in activity patterns between weekdays and weekends?**